

Batch Normalization

Pavlos Protopapas

Feature Normalization

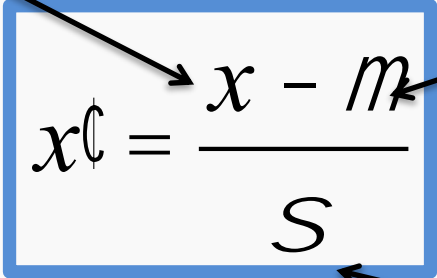
Good practice to **normalize** features before applying learning algorithm:

Feature vector

$$x_{\text{norm}} = \frac{x - m}{s}$$

Vector of mean feature values

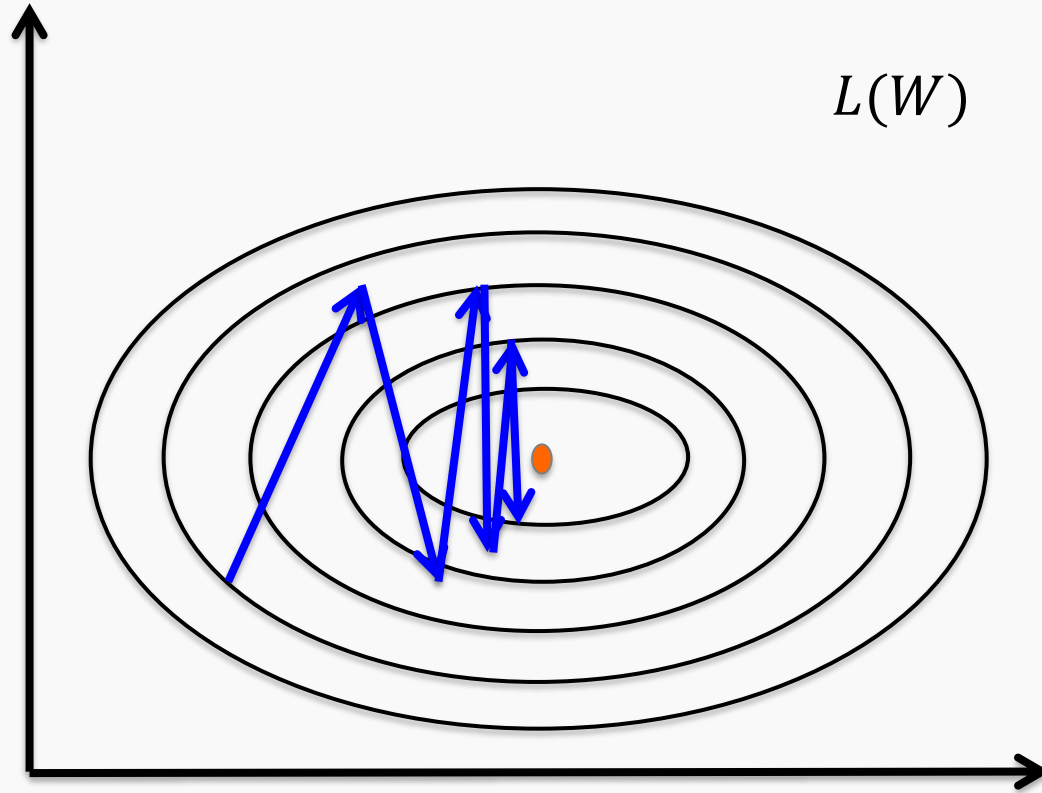
Vector of SD of feature values

A blue rectangular box contains the formula $x_{\text{norm}} = \frac{x - m}{s}$. Three arrows point from text labels to parts of the formula: 'Feature vector' points to x , 'Vector of mean feature values' points to m , and 'Vector of SD of feature values' points to s .


Features in **same scale**: mean 0 and variance 1

Feature Normalization

Speeds up learning



Before normalization

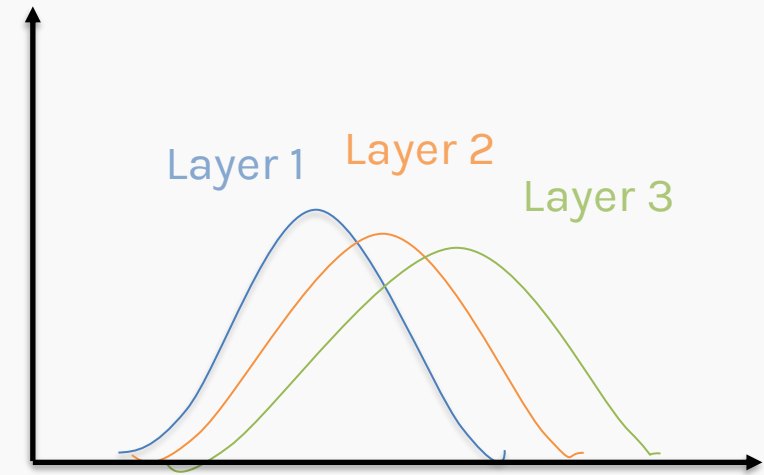


After normalization

Note: This is an ideal case scenario. In real, loss landscapes are much more complex.

Internal Covariance Shift


Even if you normalize your data to mean 0 and variance 1, the shape of the distribution may still change as you propagate through the layers of your neural network.



Distribution of the outputs of different layers

Internal Covariance Shift

Each hidden layer changes distribution of inputs to next layer: *slows down learning*



Internal Covariance Shift

We know that:

$$\frac{dL}{dW^{(1)}} = \frac{dL}{dH^{(n)}} \times \frac{dH^{(n-1)}}{dH^{(n-2)}} \times \dots \times \frac{dH^{(2)}}{dH^{(1)}} \times \frac{dH^{(1)}}{dW^{(1)}}$$

Internal Covariance Shift


We know that:

$$\frac{dL}{dW^{(1)}} = \frac{dL}{dH^{(n)}} \times \frac{dH^{(n-1)}}{dH^{(n-2)}} \times \dots \times \frac{dH^{(2)}}{dH^{(1)}} \times \frac{dH^{(1)}}{dW^{(1)}}$$




Let's look at this gradient.

Internal Covariance Shift

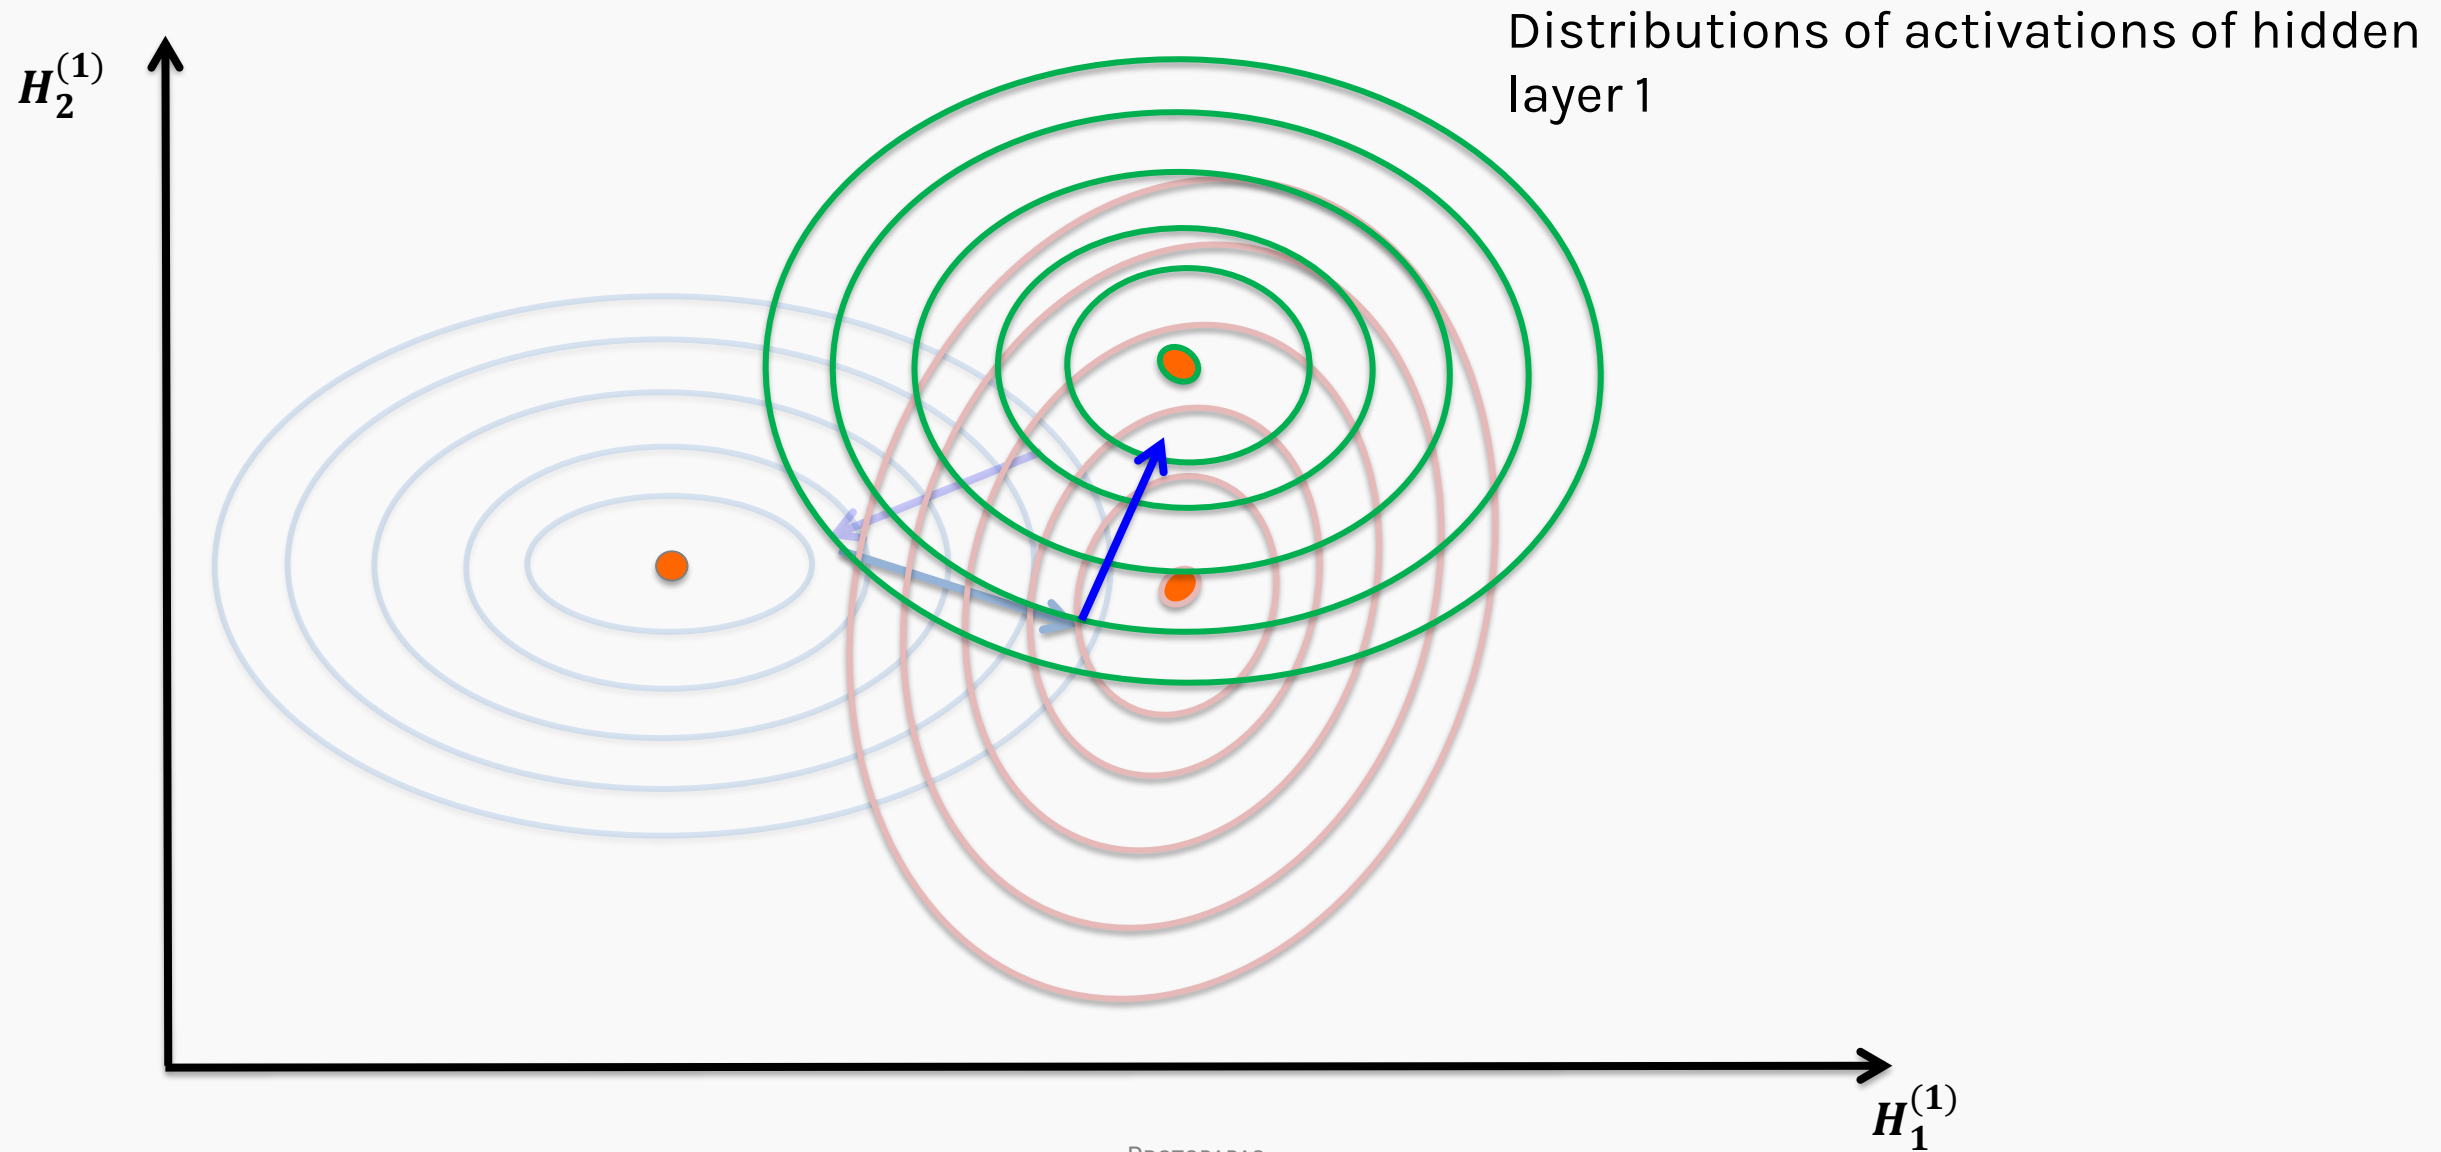


Distributions of activations of hidden layer 1

Internal Covariance Shift



Internal Covariance Shift

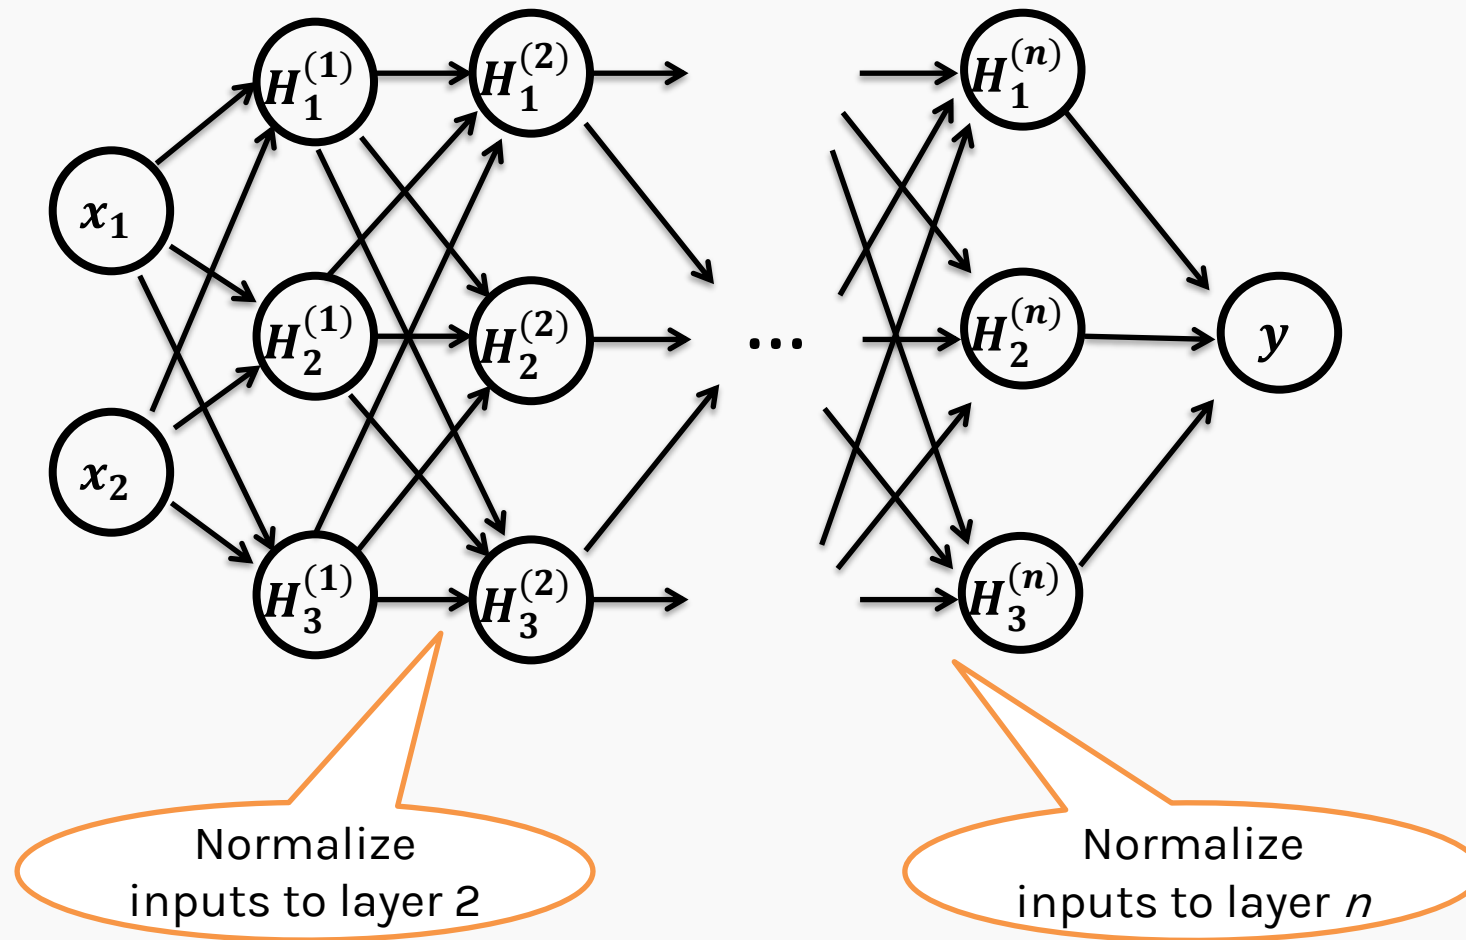




So, How can this problem be solved?

Internal Covariance Shift Solution

We normalize inputs to every hidden layer.



Batch Normalization

Training time:

Batch of activations for a given layer to normalize

For a given
hidden layer

$$H = \begin{bmatrix} H_{11} & \cdots & H_{1K} \\ \vdots & \ddots & \vdots \\ H_{N1} & \cdots & H_{NK} \end{bmatrix}$$

N data points
in batch

K hidden units
activations

Batch Normalization

Training time:

Batch of activations for a given layer to normalize

$$H = \begin{bmatrix} H_{11} & \cdots & H_{1K} \\ \vdots & \ddots & \vdots \\ H_{N1} & \cdots & H_{NK} \end{bmatrix}$$

$$H'_{ik} = \frac{H_{ik} - \mu_k}{\sigma_k}$$

Batch Normalization

Training time:

Batch of activations for a given layer to normalize

$$H = \begin{bmatrix} H_{11} & \cdots & H_{1K} \\ \vdots & \ddots & \vdots \\ H_{N1} & \cdots & H_{NK} \end{bmatrix}$$

$$H'_{ik} = \frac{H_{ik} - \mu_k}{\sigma_k}$$

$$\mu_k = \frac{1}{N} \sum_i H_{ik}$$

Mean activations across batch for node k.

Batch Normalization

Training time:

Batch of activations for a given layer to normalize

$$H = \begin{bmatrix} H_{11} & \cdots & H_{1K} \\ \vdots & \ddots & \vdots \\ H_{N1} & \cdots & H_{NK} \end{bmatrix}$$

$$H'_{ik} = \frac{H_{ik} - \mu_k}{\sigma_k}$$

$$\mu_k = \frac{1}{N} \sum_i H_{ik}$$


Mean activation across batch for node

$$\sigma_k = \sqrt{\frac{1}{N} \sum_i (H_{ik} - \mu_k)^2 + \delta}$$

SD of each unit across batch

When calculating the variance, we add a small constant to the variance to prevent potential divisions by zero.

Batch Normalization




Where,

m : Number of training examples in the batch

$H_i^{(1)}$: Hidden layer activation of the first hidden layer for the i^{th} training example

δ : A small constant

Batch Normalization



Where,

m : Number of training examples in the batch

$H_i^{(2)}$: Hidden layer activation of the second hidden layer for the i^{th} training example

δ : A small constant

Batch Normalization

Training time:

- Normalization can reduce expressive power
- Instead use:

$$H'_{ik} = \gamma H'_{ik} + \beta$$

Batch Normalization

Training time:

- Normalization can reduce expressive power
- Instead use:

$$H'_{ik} = \gamma H'_{ik} + \beta$$

Learnable parameters



When do we normalize: before or after activation?

Before activation

Batch Normalization

We have the equation

$$h^{(2)} = W a^{(1)} + b$$

where

$a^{(1)}$: Activation of the first hidden layer

$h^{(2)}$: the output of the second hidden layer
w/o activation

If we do batch normalization **after** activation:

The shape of the distribution of $a^{(1)}$ is likely to change during training and limiting its mean and standard deviation will not eliminate covariate shift.

Batch Normalization

We have the equation

$$h^{(2)} = W a^{(1)} + b$$

where

$a^{(1)}$: Activation of the first hidden layer

$h^{(2)}$: the output of the second hidden layer
w/o activation

If we do batch normalization **before** activation:

$W a^{(1)} + b$ is very likely to have a symmetric, non-sparse distribution;
normalizing it is likely to produce activations with a stable distribution.



We saw how batch normalization works during training, but what about **prediction**, when we might not have a complete batch!

Evaluation

Evaluation time:

- Calculate the running average of the mean and standard deviation.
- For every batch:

Decay
parameter

Use this for
evaluation

$$\mu_{global} = \alpha \mu_{global} + (1 - \alpha) \mu_k$$

$$\sigma_{global} = \alpha \sigma_{global} + (1 - \alpha) \sigma_k$$

Batch Normalization

Evaluation time:

Hidden activations will be a vector as there are no batches.

$$H = [H_1 \quad \dots \quad H_K]$$

Batch Normalization

Evaluation time:

Use the global statistics to normalize the node activations.

$$H = [H_1 \quad \dots \quad H_K]$$

For each hidden node k :

$$H'_k = \frac{H_k - \mu_{global}}{\sigma_{global}}$$

← Estimated global mean of each unit activation.

↑
Estimated global SD of each unit activation.